

Chapter 06

Game Playing

Instructor

LE Thanh Sach, Ph.D.

Instructor's Information

LE Thanh Sach, Ph.D.

Office:

Department of Computer Science,
Faculty of Computer Science and Engineering,
HoChiMinh City University of Technology.

Office Address:

268 LyThuongKiet Str., Dist. 10, HoChiMinh City,
Vietnam.

E-mail: LTSACH@cse.hcmut.edu.vn

E-home: <http://cse.hcmut.edu.vn/~ltsach/>

Tel: (+84) 83-864-7256 (Ext: 5839)

Acknowledgment

The slides in this PPT file are composed using the materials supplied by

✍ **Prof. Stuart Russell and Peter Norvig:** They are currently from University of California, Berkeley. They are also the author of the book “Artificial Intelligence: A Modern Approach”, which is used as the textbook for the course

✍ **Prof. Tom Lenaerts,** from Université Libre de Bruxelles

Outline

- ❖ What are games?
- ❖ Optimal decisions in games
 - ✂ Which strategy leads to success?
- ❖ α - β pruning
- ❖ Games of imperfect information

What are and why study games?

- ❖ Games are a form of *multi-agent environment*
 - ✂ What do other agents do and how do they affect our success?
 - ✂ Cooperative vs. competitive multi-agent environments.
 - ✂ Competitive multi-agent environments give rise to adversarial problems a.k.a. *games*
- ❖ Why study games?
 - ✂ Fun; historically entertaining
 - ✂ Interesting subject of study because they are hard
 - ✓ Chess game:
average branch factor: 35, each player: 50 moves → Search tree: 35^{100} nodes ☹

Relation of Games to Search

❖ Search – no adversary

- ✍ Solution is (heuristic) method for finding goal
- ✍ Heuristics and CSP techniques can find *optimal* solution
- ✍ Evaluation function: estimate of cost from start to goal through given node
- ✍ Examples: path planning, scheduling activities

❖ Games – adversary

- ✍ Solution is strategy (strategy specifies move for every possible opponent reply).
- ✍ Time limits force an *approximate* solution
- ✍ Evaluation function: evaluate “goodness” of game position
- ✍ Examples: chess, checkers, Othello, backgammon


Types of Games

	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information		bridge, poker, scrabble nuclear war

Game setup

- ❖ Two players: MAX and MIN
- ❖ MAX moves first and they take turns until the game is over. Winner gets award, loser gets penalty.
- ❖ Games as search:
 - ✍ Initial state: e.g. board configuration of chess
 - ✍ Successor function: list of (move,state) pairs specifying legal moves.
 - ✍ Terminal test: Is the game finished?
 - ✍ Utility function: Gives numerical value of terminal states.
E.g. win (+1), loose (-1) and draw (0) in tic-tac-toe (next)
- ❖ MAX uses search tree to determine next move.

Partial Game Tree for Tic-Tac-Toe



Optimal strategies

- ❖ Find the contingent *strategy* for MAX assuming an infallible MIN opponent.
- ❖ Assumption: Both players play optimally !!
- ❖ Given a game tree, the optimal strategy can be determined by using the minimax value of each node:

MINIMAX-VALUE(n)=

UTILITY(n)

If n is a terminal


$\max_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$

If n is a max node


$\min_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$

If n is a min node


Two-Ply Game Tree




Two-Ply Game Tree



Two-Ply Game Tree



Two-Ply Game Tree

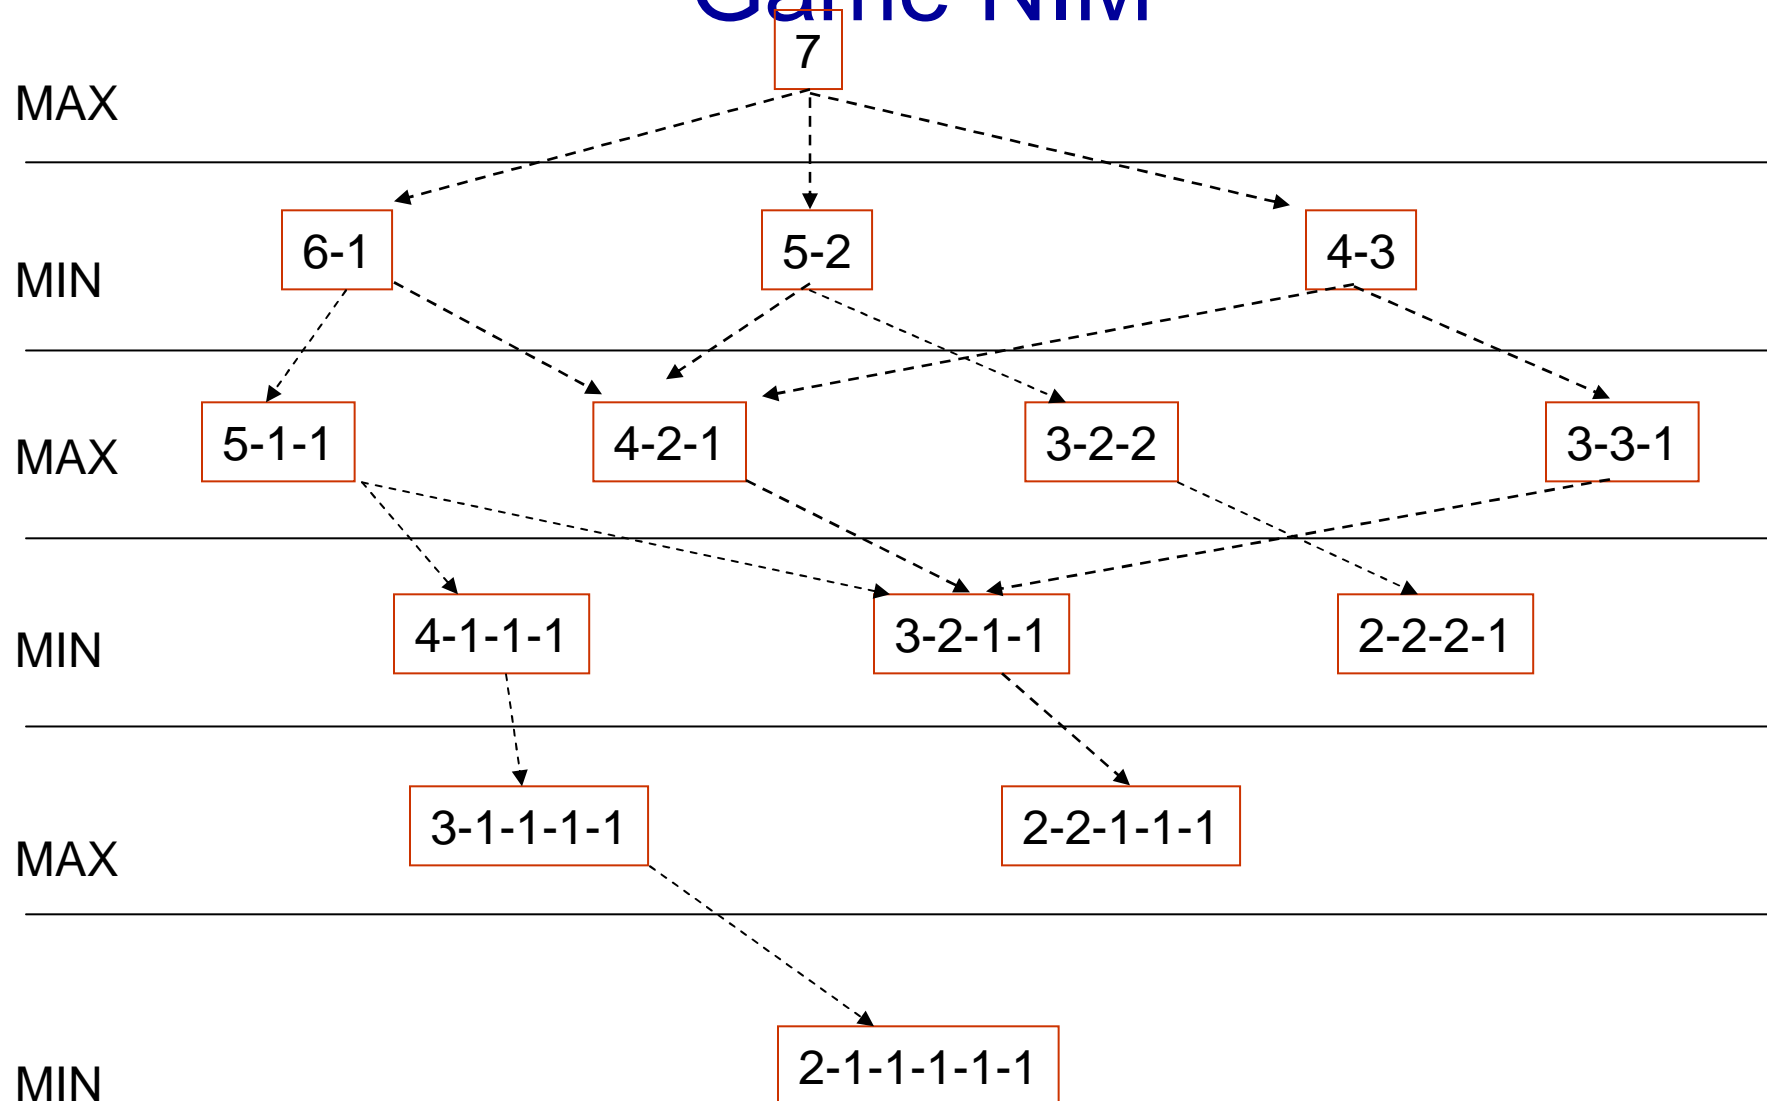


Minimax maximizes the worst-case outcome for max.

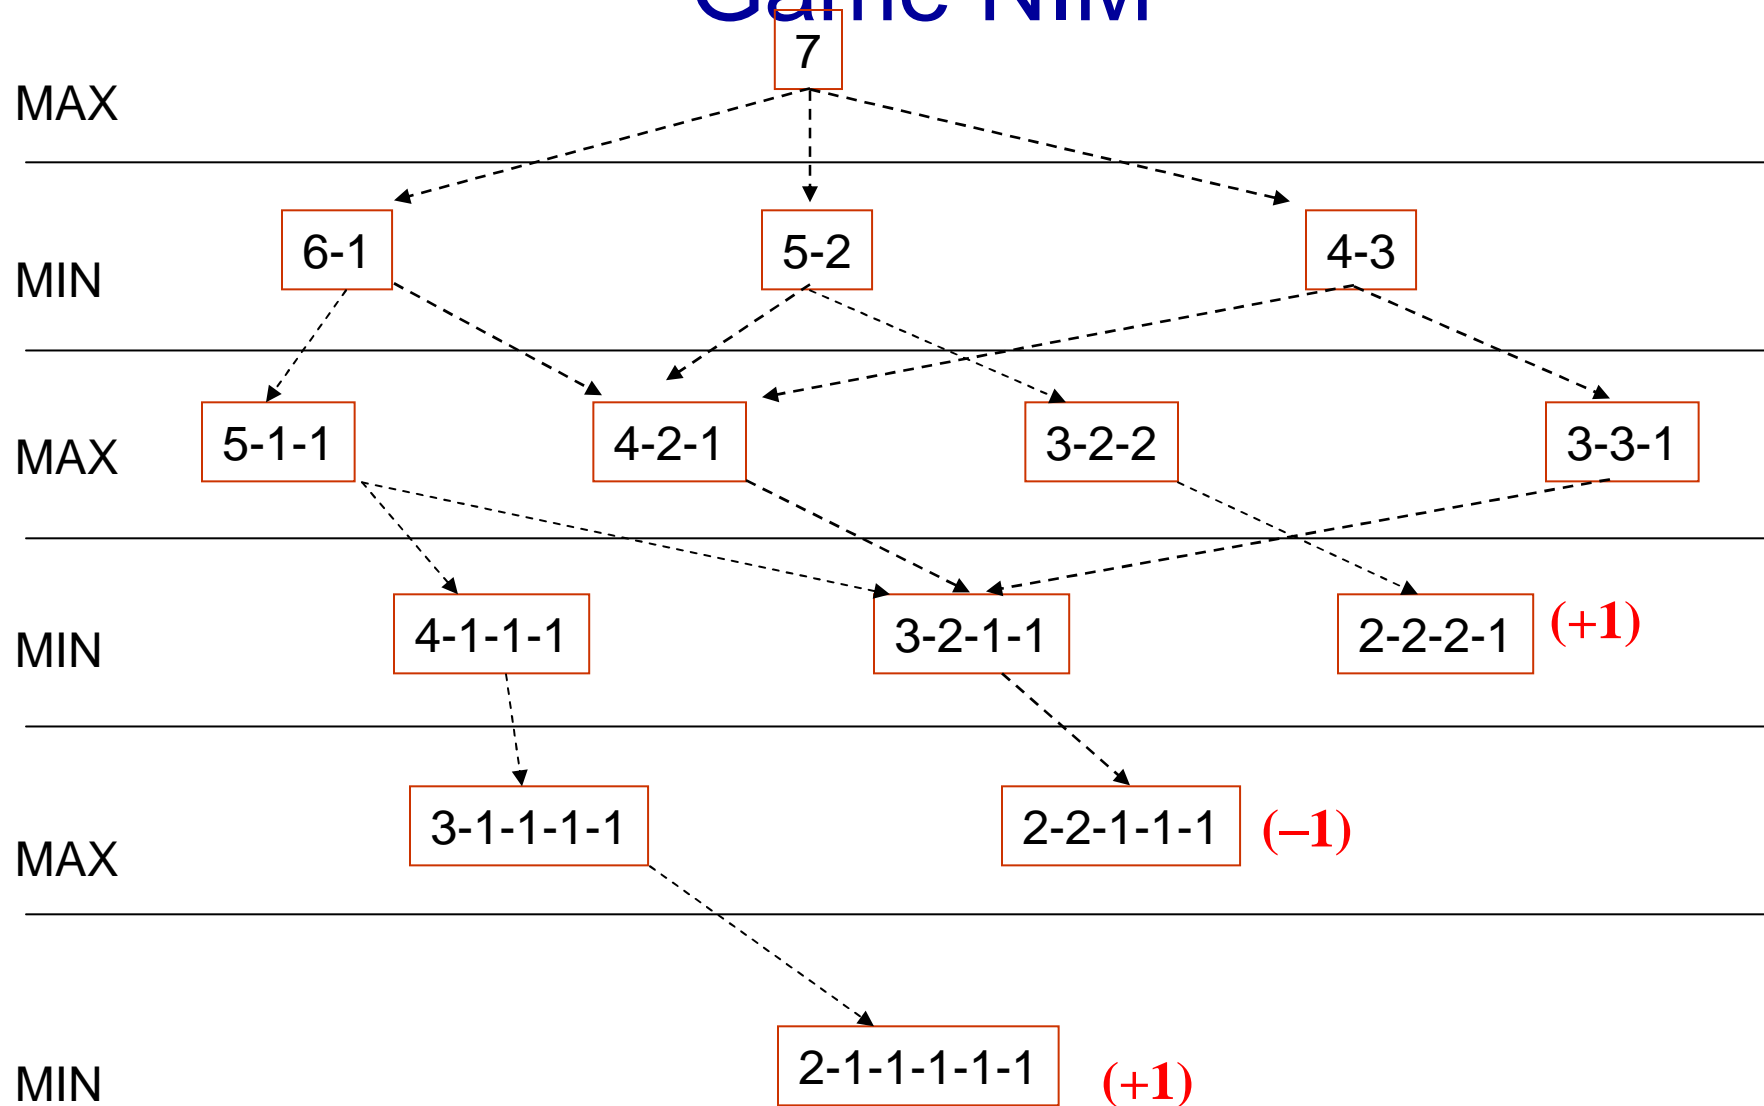
Game NIM

- ❖ Given N rods in n ($n \geq 1$) heaps, two players, take turns
- ❖ At each turn, player has to select one heap and separate it into two different heaps having different number of rods
- ❖ A player loses the game if he/she can not select a heap for separation.

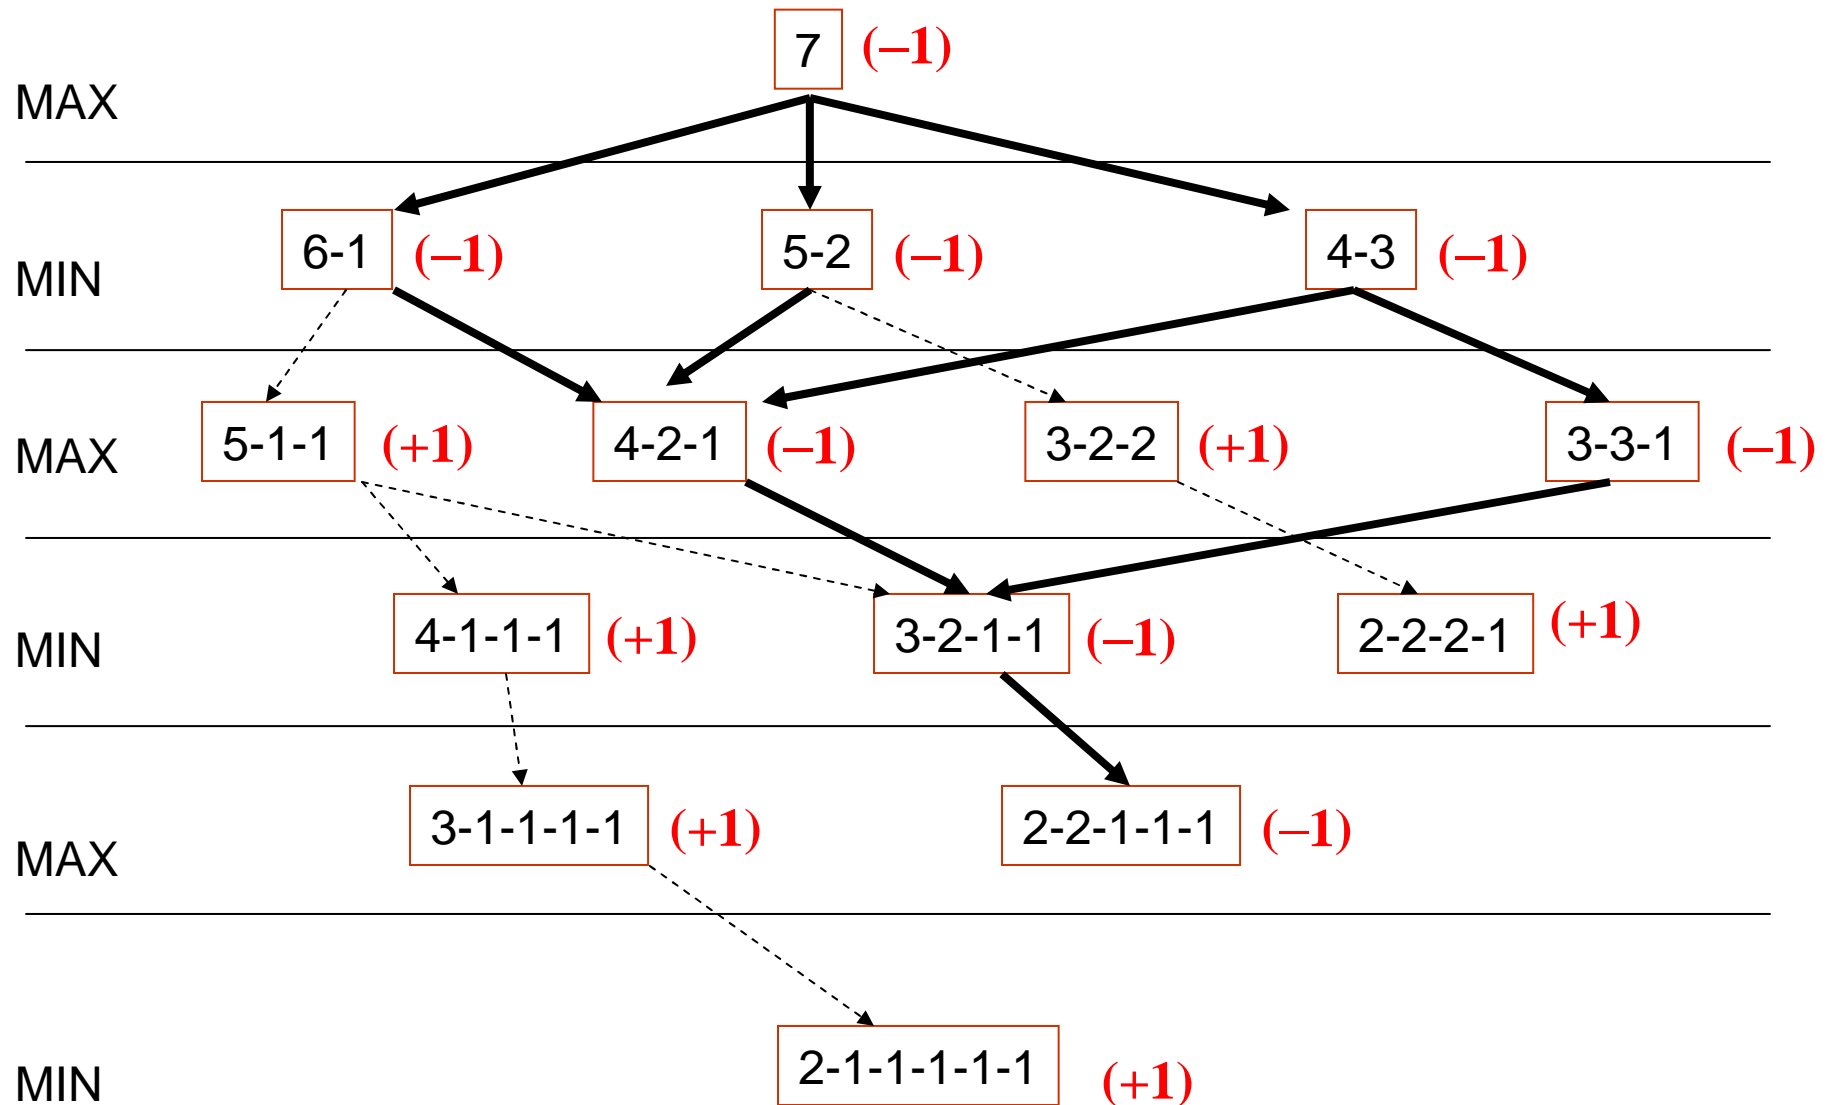
Game NIM



Game NIM



Game NIM



What if MIN does not play optimally?

- ❖ Definition of optimal play for MAX assumes MIN plays optimally: maximizes worst-case outcome for MAX.
- ❖ But if MIN does not play optimally, MAX will do even better. [Can be proved.]

Minimax Algorithm

function MINIMAX-DECISION(*state*) **returns** *an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a,s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)





$v \leftarrow +\infty$

for *a,s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

Properties of Minimax

Criterion	Minimax
Complete?	Yes 
Time	$O(b^m)$ 
Space	$O(bm)$ 
Optimal?	Yes 

Multiplayer games

- ❖ Games allow more than two players
- ❖ Solution:
 - ✎ Single minimax values become vectors
 - ✎ For example, 3 players: A, B, C
 - ✎ Values at each node: $[V_A, V_B, V_C]$
 - ✓ V_A : measure the utility of the state under in viewpoint of A
 - ✓ V_B : measure the utility of the state under in viewpoint of B
 - ✓ V_C : measure the utility of the state under in viewpoint of C
- ❖ Multiplayer → Making Alliances

Multiplayer games


to move

A

B

C

A



Multiplayer games


- ❖ Multiplayer → Strategy: **Making Alliances**, look like “Three Kingdoms” in Chinese history
 - ✍ A and B are in their weak points
 - ✍ C is in its stronger point
 - ✍ → It is a reasonable that A and B make a alliance and together attack C
 - ✍ ...
 - ✍ However, as soon as C weakens the alliance is break down.

Problem of minimax search


- ❖ Number of games states is exponential to the number of moves.
 - ✍ **Solution: Do not examine every node**
 - ✍ **==> Alpha-beta pruning**
 - ✓ Alpha = value of best choice found so far at any choice point along the MAX path
 - ✓ Beta = value of best choice found so far at any choice point along the MIN path
- ❖ Revisit example ...

Alpha-Beta Example


Do DF-search until first leaf




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)



Alpha-Beta Example (continued)



Alpha-Beta Example (continued)



Alpha-Beta Algorithm

function ALPHA-BETA-SEARCH(*state*) **returns** *an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$


return v

Alpha-Beta Algorithm



```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for a,s in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

General alpha-beta pruning


- ❖ Consider a node n somewhere in the tree
- ❖ If player has a better choice at
 - ✂ **Parent node of n**
 - ✂ **Or any choice point further up**
- ❖ n will **never** be reached in actual play.
- ❖ Hence when enough is known about n , it can be pruned.



Alpha-Beta Pruning




Alpha-Beta Example



Suppose that the search tree has the form as shown in the figure

Alpha-Beta Example



Final Comments about Alpha-Beta Pruning

- ❖ Pruning does not affect final results
- ❖ Entire subtrees can be pruned.
- ❖ Good move *ordering* improves effectiveness of pruning
- ❖ With “perfect ordering,” time complexity is $O(b^{m/2})$
 - ✂ Branching factor of \sqrt{b} !!
 - ✂ Alpha-beta pruning can look twice as far as minimax in the same amount of time
- ❖ Repeated states are again possible.
 - ✂ Store them in memory = transposition table

Games of imperfect information

- ❖ Minimax and alpha-beta pruning require too much leaf-node evaluations.
- ❖ May be impractical within a reasonable amount of time.
- ❖ SHANNON (1950):
 - ✂ Cut off search earlier (replace TERMINAL-TEST by CUTOFF-TEST)
 - ✂ Apply heuristic evaluation function EVAL (replacing utility function of alpha-beta)

Cutting off search

- ❖ Change:

- ~~✂~~ **if** `TERMINAL-TEST(state)` **then return** `UTILITY(state)`

into

- ~~✂~~ **if** `CUTOFF-TEST(state,depth)` **then return** `EVAL(state)`

- ❖ Introduces a fixed-depth limit *depth*


- ~~✂~~ Is selected so that the amount of time will not exceed what the rules of the game allow.

- ❖ When cutoff occurs, the evaluation is performed.


Heuristic EVAL

- ❖ Idea: produce an estimate of the expected utility of the game from a given position.
- ❖ Performance depends on quality of EVAL.
- ❖ Requirements:
 - ✂ EVAL should order terminal-nodes in the same way as UTILITY.
 - ✂ Computation may not take too long.
 - ✂ For non-terminal states the EVAL should be strongly correlated with the actual chance of winning.
- ❖ Only useful for quiescent (no wild swings in value in near future) states


Heuristic EVAL example




(a) White to move
Fairly even



(b) Black to move
White slightly better




(c) White to move
Black winning




(d) Black to move
White about to lose

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$


Heuristic EVAL example




(a) White to move
Fairly even



(b) Black to move
White slightly better



(c) White to move
Black winning




(d) Black to move
White about to lose

*Addition assumes
independence*


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

Heuristic difficulties

Heuristic counts pieces won




(a) White to move



(b) White to move


Discussion

- ❖ Examine section on state-of-the-art games yourself
- ❖ Minimax assumes right tree is better than left, yet ...
 - ✂ Return probability distribution over possible values
 - ✂ Yet expensive calculation



Discussion

- ❖ Utility of node expansion
 - ✂ Only expand those nodes which lead to significantly better moves
- ❖ Both suggestions require meta-reasoning



Summary

- ❖ Games are fun (and dangerous)
- ❖ They illustrate several important points about AI
 - ✂ Perfection is unattainable -> approximation
 - ✂ Good idea what to think about
 - ✂ Uncertainty constrains the assignment of values to states
- ❖ Games are to AI as grand prix racing is to automobile design.